




Maroc Ynov Campus  
M2 – Master Développement Web et Mobile

Microservices Architecture

# Event-Driven Data Management


by  
**Driss ALLAKI**



# Event-Driven Data Management

The general goal

Implementing transactions and queries  
in a microservice architecture  
using asynchronous messaging




# Event-Driven Data Management

- Problem Statement : **Distributed Data Management** (*Transactions and Queries*)
- One solution of the **Transactions** challenge:
  - The **Saga** pattern
  - Coordination sagas
    - Choreography-based saga
    - Orchestration-based saga
- Two Solutions of the **Queries** challenge :
  - Solution 1 : The **API Composition** pattern
  - Solution 2 : The **CQRS** pattern

# Problem Statement : The problem of Distributed Data Management

## Data Management in Monolith Applications

- A **monolithic application** typically has a **single relational database**. A key benefit of using a relational database is that your application can use **ACID transactions**, which provide some important guarantees
  - **Atomicity** : Changes are made atomically (all or nothing)
  - **Consistency** : The state of the database is always consistent
  - **Isolation** : Even though transactions are executed concurrently it appears they are executed serially
  - **Durability** : Once a transaction has committed it is not undone
- As a result, the application can simply :
  - begin a transaction
  - change (insert, update, and delete) multiple rows
  - and commit the transaction.
- Another great benefit of using a relational database is that it provides **SQL**.
- You can easily write a query that **combines data from multiple tables**.




# Problem Statement : The problem of Distributed Data Management

## Data Management in a Distributed Microservices Architecture

- Data access becomes much more complex when we move to a **microservices architecture**.
  - The **data** owned by each microservice is **private** to that microservice and can only be accessed via its API (the **Database per service** pattern)  
*(If multiple services access the same data, schema updates require time-consuming, coordinated updates to all the services => Tight coupling and dependency between services)*
  - Different microservices often use **different kinds of databases** (the **polyglot persistence** approach)

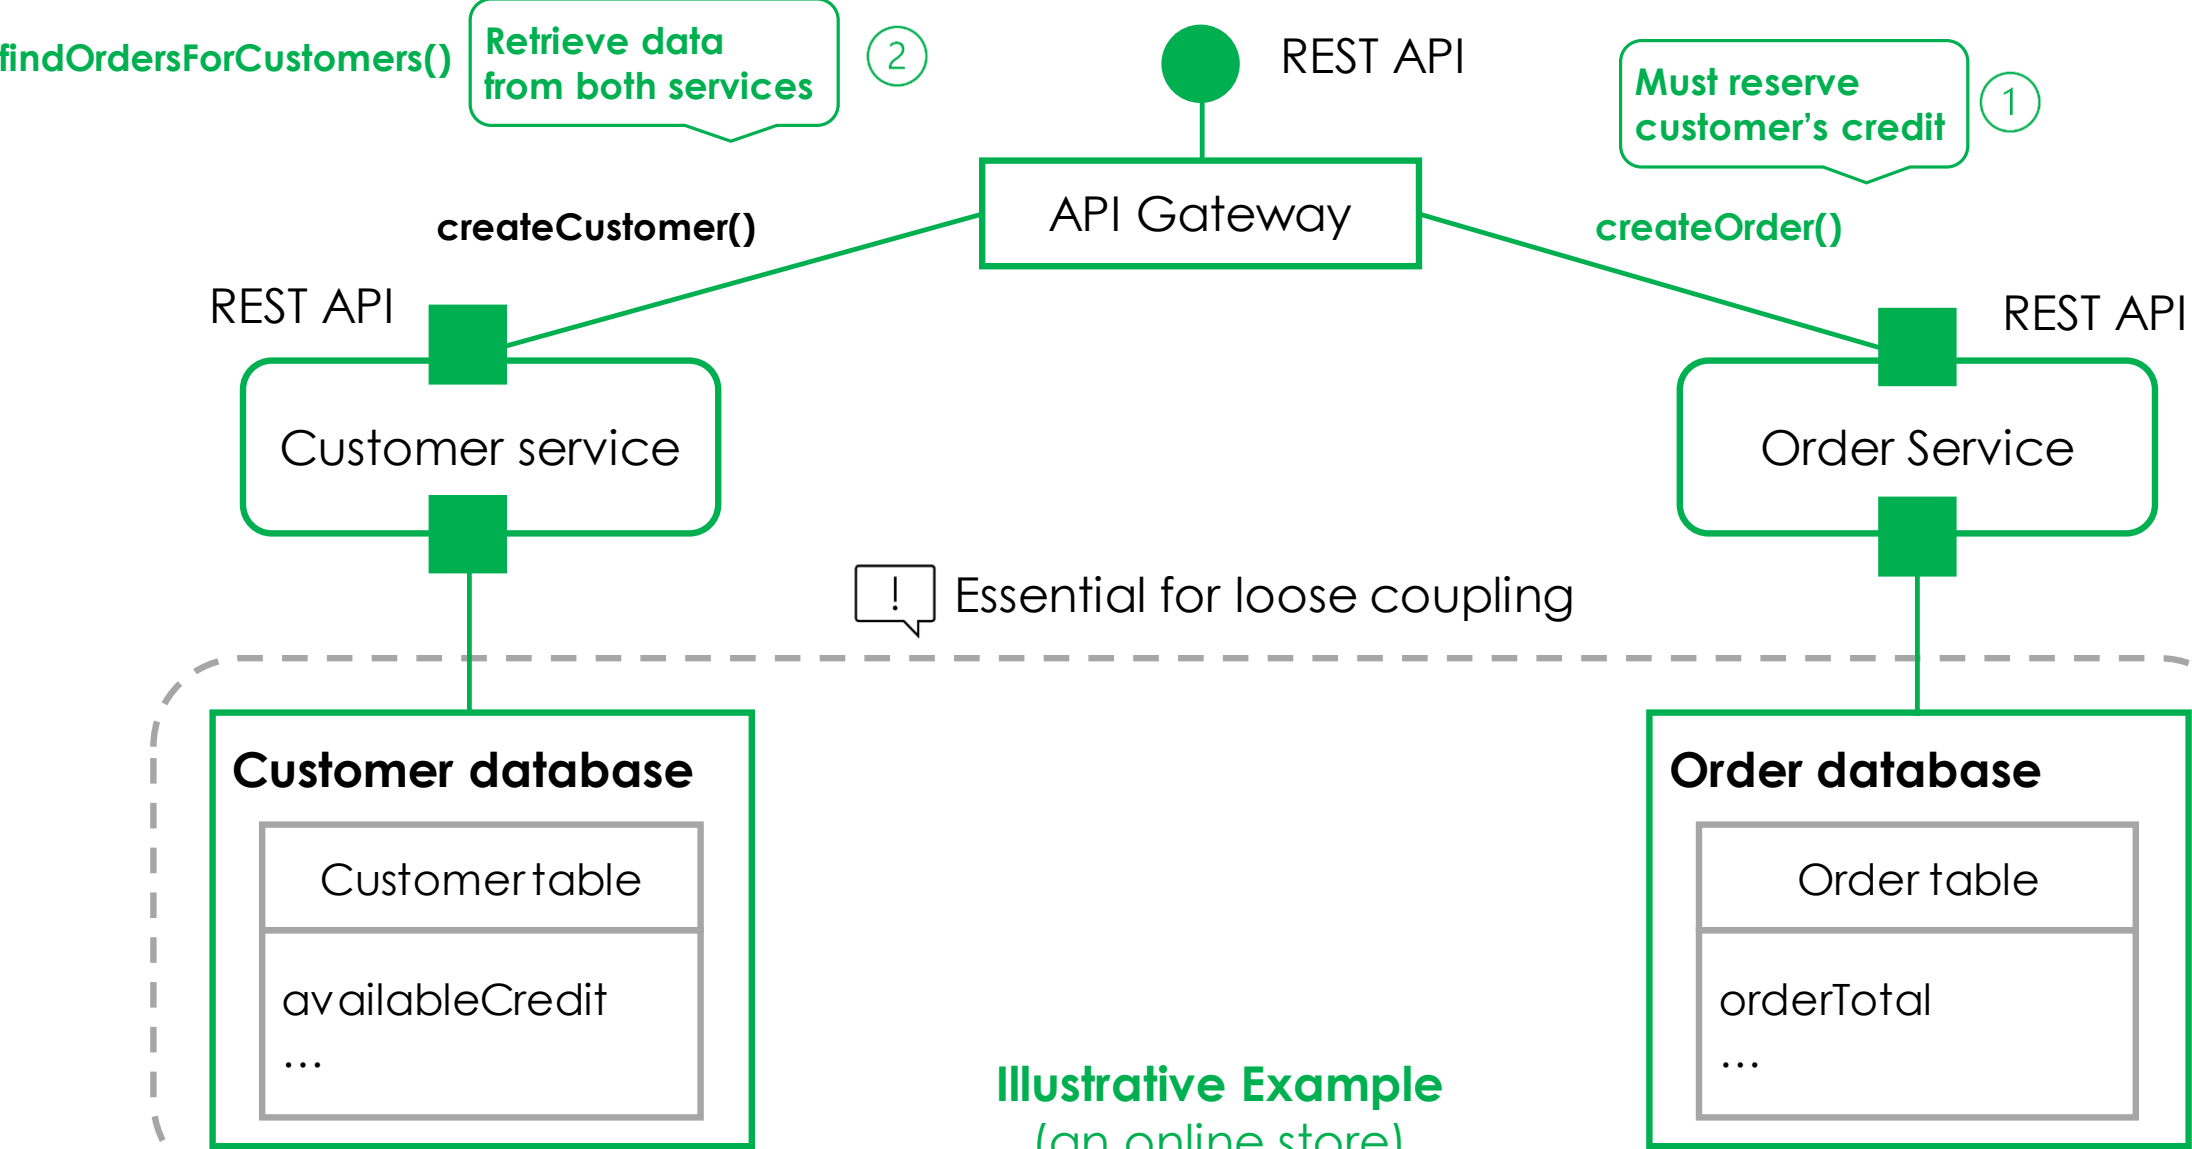
- ⊕ **A partitioned, polyglot-persistent architecture for data storage** has many benefits, including :
  - **loosely coupled services**
  - better **performance**
  - better **scalability**

- ⊖ It introduces some **distributed data management challenges**.



# Problem Statement : The problem of Distributed Data Management

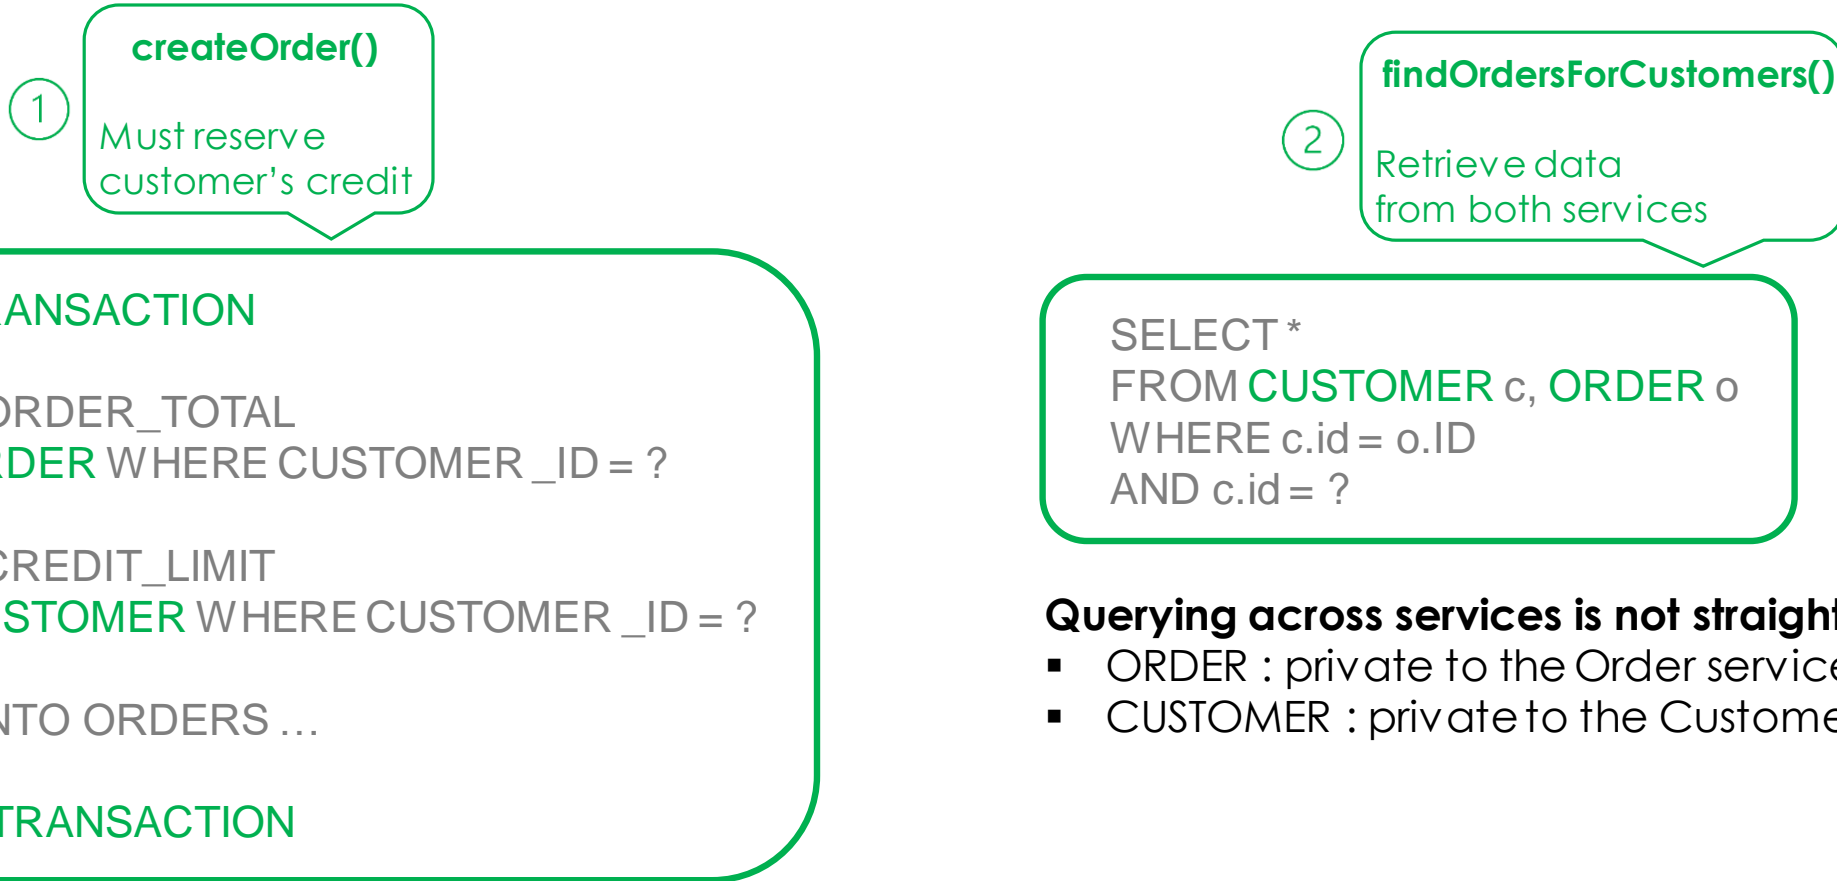
How to achieve consistency across multiple microservices



# Problem Statement : The problem of Distributed Data Management

How to achieve consistency across multiple microservices

**Illustrative Example** (an online store)



**Querying across services is not straightforward**

- ORDER : private to the Order service
- CUSTOMER : private to the Customers service

**No ACID transactions that span services**

- Distributed transaction
- ORDER : private to the Order service
- CUSTOMER : private to the Customers service

# Problem Statement : The problem of Distributed Data Management

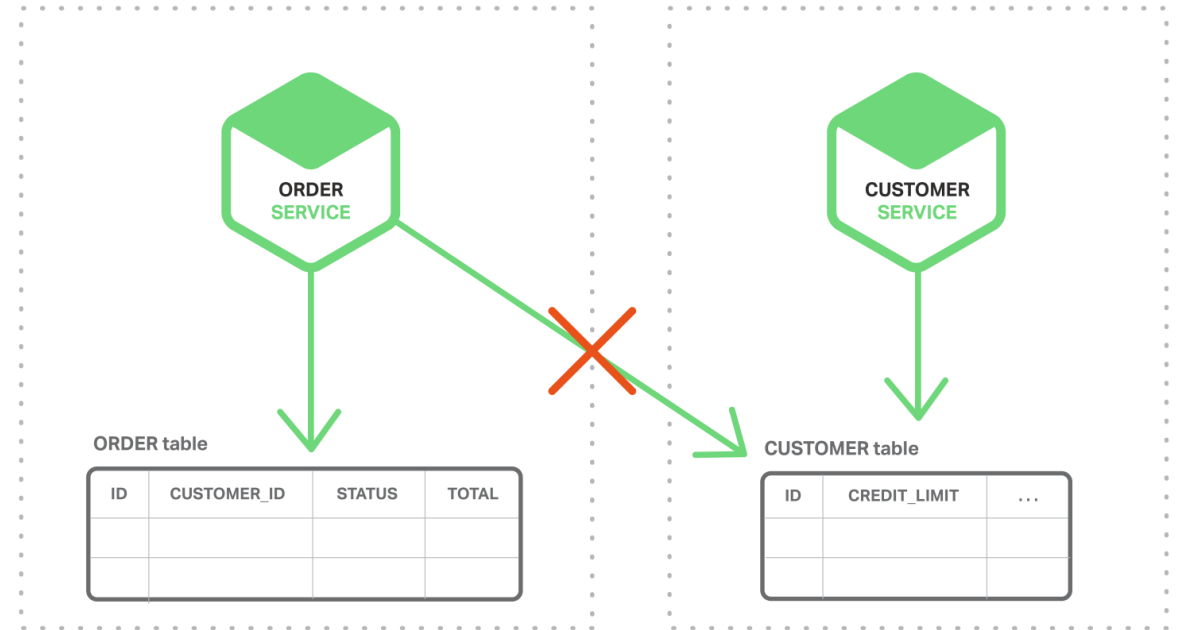
How to achieve consistency across multiple microservices

## Illustrative Example : (an online store)

- The *Customer* Service maintains information about customers, including their credit lines.
- The *Order* Service manages orders and must verify that a new order doesn't exceed the customer's credit limit.

In a microservices architecture the ORDER and CUSTOMER tables are private to their respective services.

(The Order Service cannot access the CUSTOMER table directly. It can only use the API provided by the Customer Service.)




## Challenges :

- ① How to implement **transactions** that span services?
- ② How to implement **queries** that retrieve data from multiple services ?

! Maintaining data consistency across services and databases is of utmost importance





# Event-Driven Data Management

- Problem Statement : **Distributed Data Management** (*Transactions and Queries*)
- One solution of the **Transactions** challenge:
  - The **Saga** pattern
  - Coordination sagas
    - Choreography-based saga
    - Orchestration-based saga
- Two Solutions of the **Queries** challenge :
  - Solution 1 : The **API Composition** pattern
  - Solution 2 : The **CQRS** pattern

# Solution : Managing transactions with SAGAs

Use SAGAs pattern instead of 2PC

2PC is not an option for Microservices

Read more about  
[2PC \(Two-Phase Commit protocol\)](#)

Distributed Transaction (2PC)

Order  
Service

Customer  
Service



Saga

Order  
Service

Local  
Transaction

Message/event

Customer  
Service

Local  
Transaction

Message/event

Order  
Service

Local  
Transaction

**SAGAS**

*Hector Garcia-Molina  
Kenneth Salem*

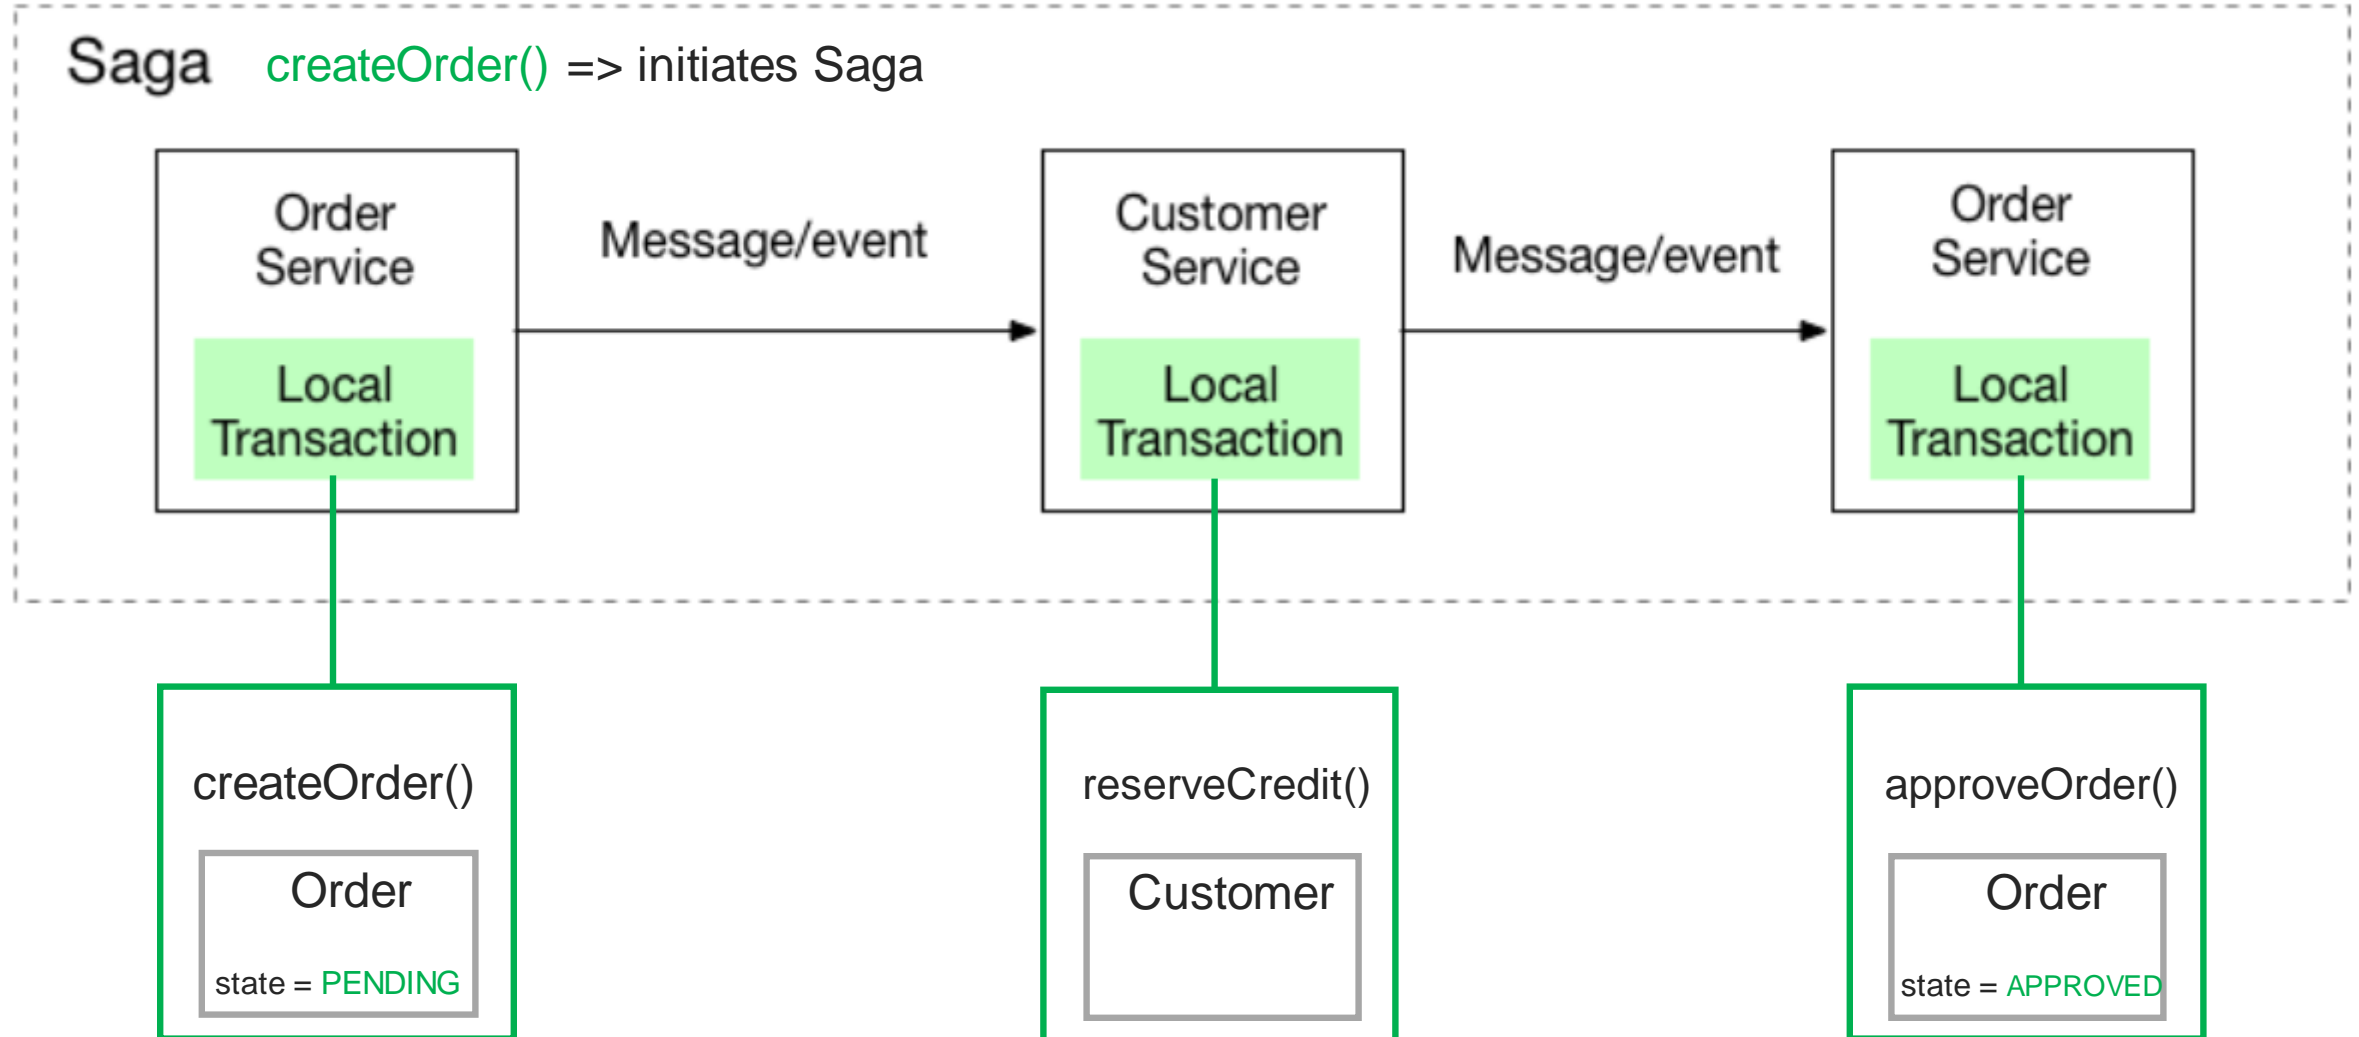
**Department of Computer Science  
Princeton University  
Princeton, N J 08544**

*From a 1987 paper*

- A **saga** is a **sequence of local transactions** (one saga step = a transaction local to a service)
- Each local transaction **updates the database** and **publishes a message or an event** to trigger the next local transaction in the saga.

# Solution : Managing transactions with SAGAs

## “Create Order” SAGA example



# Solution : Managing transactions with SAGAs

Few questions to ask about Sagas

- **Question 1** : How do the saga participants communicate?
- **Answer 1** : Collaboration using **asynchronous, broker-based messaging**

- **Question 2** :

After the completion of a transaction **T(i)** “something” must decide : **what step to execute next ?**

- **Success**: which  $T(i+1)$  - branching
- **Failure**:  $C(i - 1)$

- **Answer 2** : There are two ways of coordination sagas :

## Choreography

(Distributed decision making)

Each local transaction publishes domain events that trigger local transactions in other services

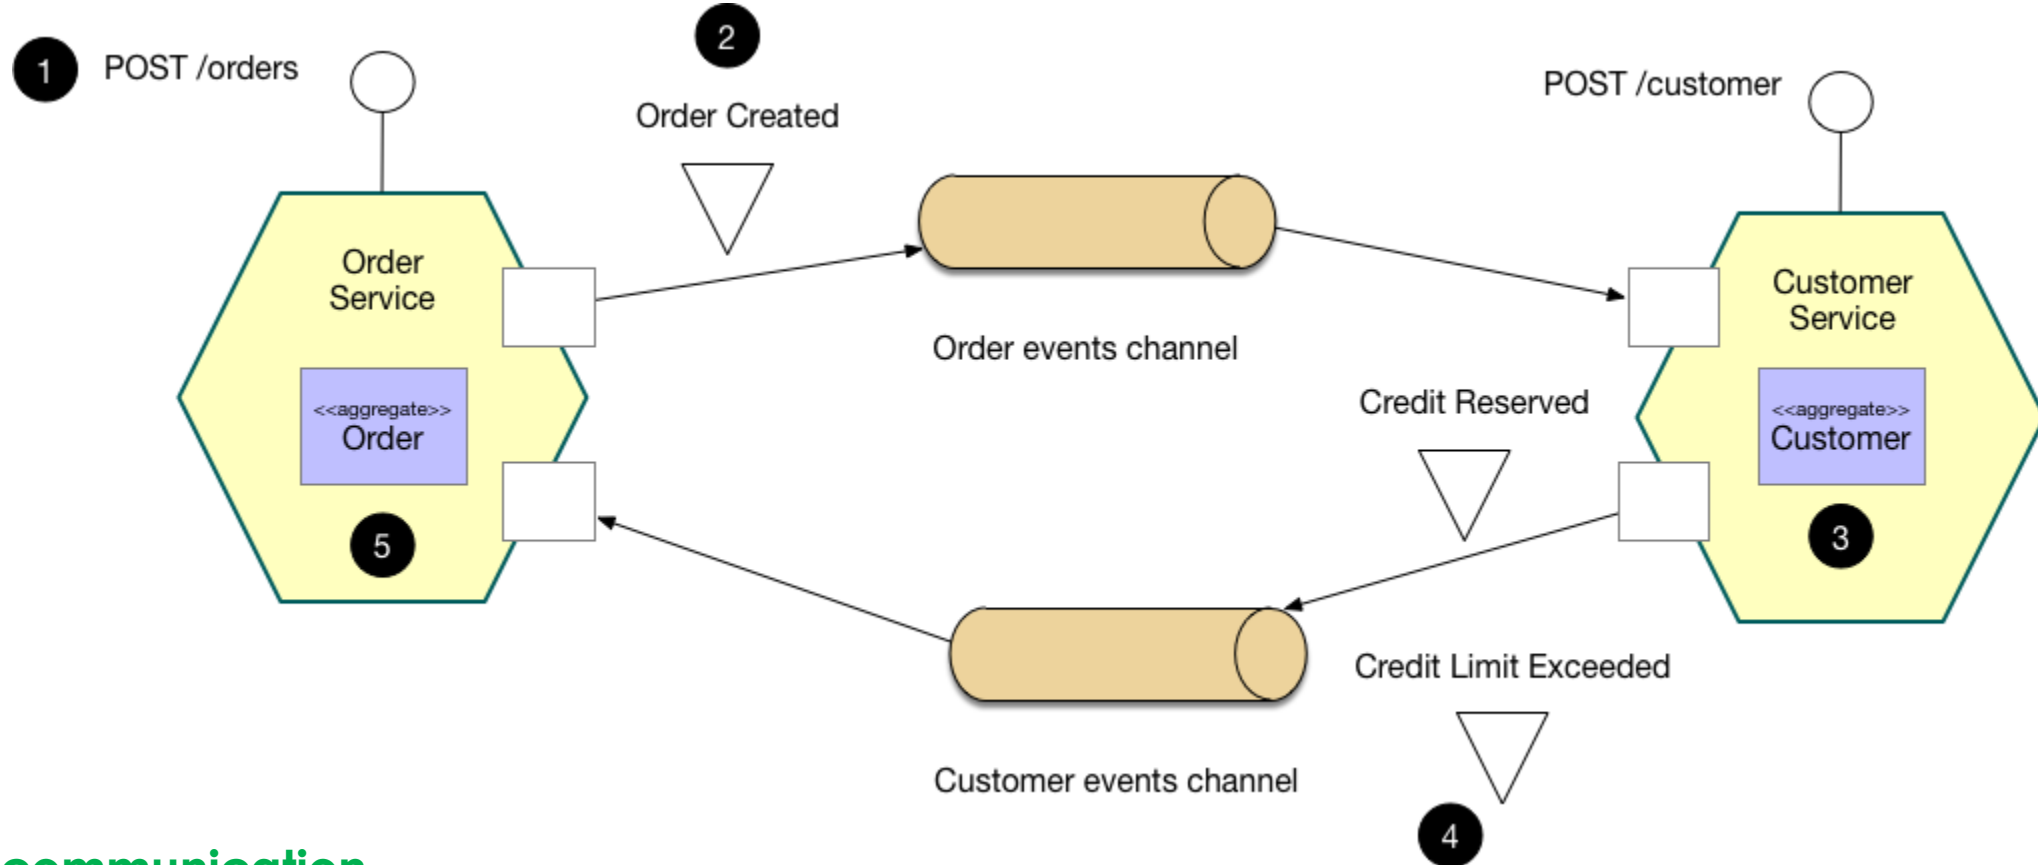
## Orchestration

(Centralized decision making)

An orchestrator (object) tells the participants what local transactions to execute

# Solution : Managing transactions with SAGAs

## Choreography-based saga : Example






## Pub/Sub communication


1. The **Order Service** receives the `POST /orders` request and creates an **Order** in a **PENDING** state
2. It then emits an `Order Created` event
3. The **Customer Service**'s event handler attempts to reserve credit
4. It then emits an event indicating the outcome
5. The **OrderService**'s event handler either approves or rejects the **Order**





# Solution : Managing transactions with SAGAs

## Choreography-based saga : Benefits and Drawbacks



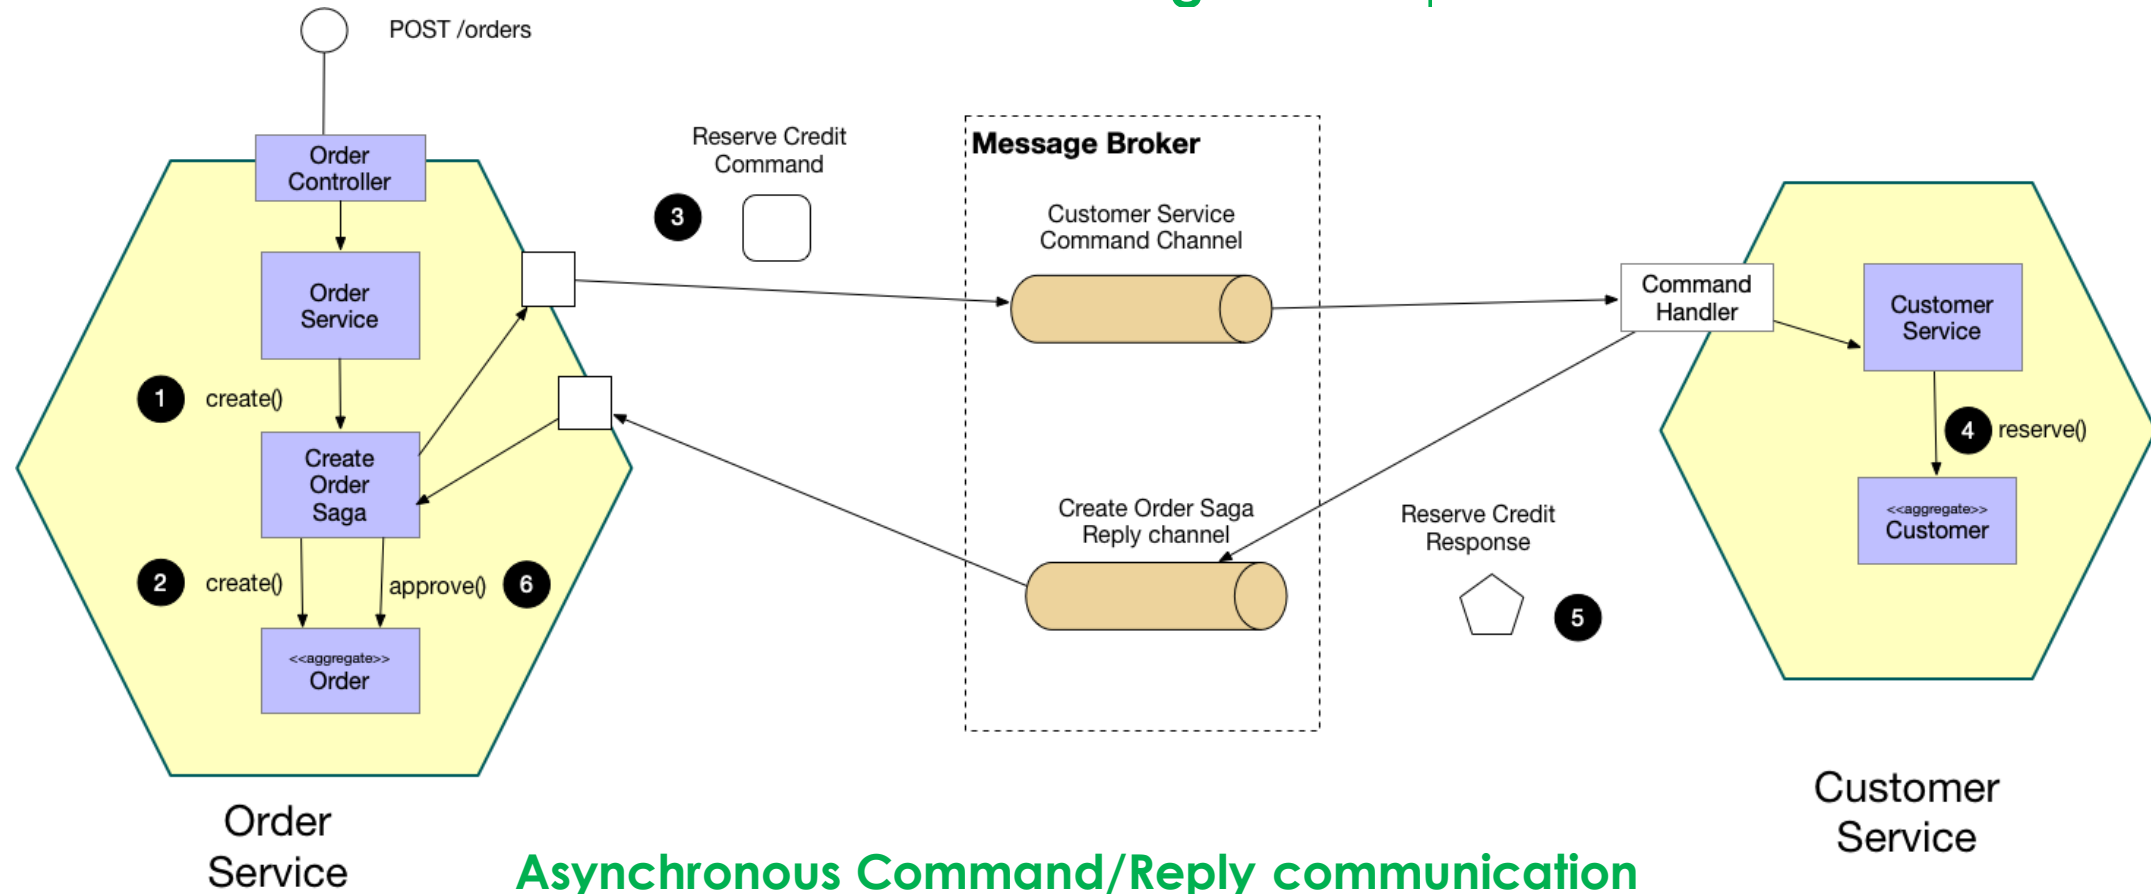
-  Simple, especially when using event sourcing
-  Participants are loosely coupled



-  Decentralized implementation - potentially difficult to understand
-  Cyclic dependencies - services listen to each other's events,  
*e.g. Customer Service must know about all Order events that affect credit*
-  Overloads domain objects,  
*e.g. Order and Customer know too much*
-  Events = indirect way to make something happen

# Solution : Managing transactions with SAGAs

## Orchestration-based saga : Example




1. The **Order Service** receives the `POST / orders` request and creates the **Create Order** saga orchestrator
2. The saga orchestrator creates an Order in the **PENDING** state
3. It then sends a **ReserveCredit** command to the **CustomerService**
4. The **CustomerService** attempts to reserve credit
5. It then sends back a reply message indicating the outcome
6. The saga orchestrator either approves or rejects the **Order**

# Solution : Managing transactions with SAGAs

## Orchestrator-based saga : Benefits and Drawbacks

A saga “orchestrator” is a **persistent object** that implements a state machine and **invokes** the participants




Centralized coordination logic is easier to understand



Reduced coupling,  
*e.g. Customer Service knows less. Simply has API for managing available credit.*




Reduces cyclic dependencies



Risk of smart sagas directing dumb services





# Event-Driven Data Management

- Problem Statement : **Distributed Data Management** (*Transactions and Queries*)
- One solution of the **Transactions** challenge:
  - The **Saga** pattern
  - Coordination sagas
    - Choreography-based saga
    - Orchestration-based saga
- Two Solutions of the **Queries** challenge :
  - Solution 1 : The **API Composition** pattern
  - Solution 2 : The **CQRS** pattern

# Solution (a) of challenge 2 : API composition pattern

**Challenge :** How to implement **queries** in a microservice architecture ?

## **Solution :**

Implement a query by defining an **API Composer**, which invoke the services that own the data and performs an **in-memory join** of the results.

## **Example :**

An **API Gateway** often does API composition.

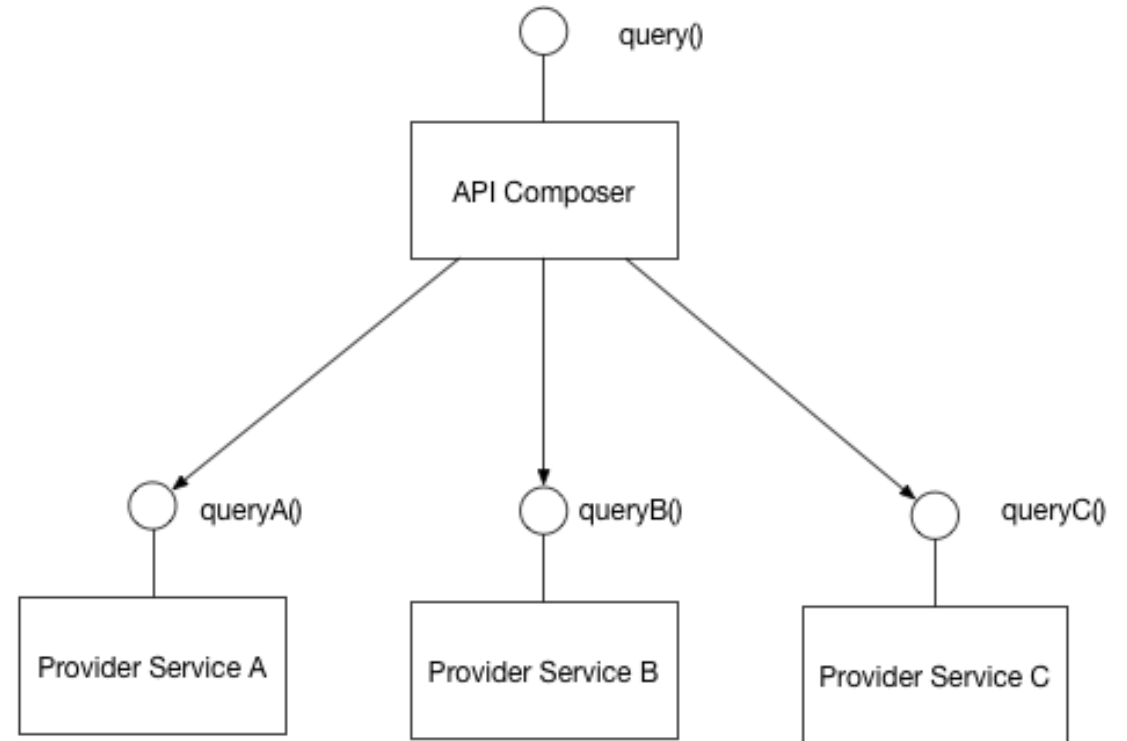
## **Resulting context :**

This pattern has the following benefit:

- ⊕ *It's a simple way to query data in a microservice architecture*

This pattern has the following drawback:

- ⊖ *Some queries would result inefficient (in-memory joins of large datasets)*



# Solution (a) of challenge 2 : API composition pattern

## Example of an API composition limitation

Find recent, valuable customers

```
SELECT *  
FROM CUSTOMER c, ORDER o  
WHERE c.id = o.ID  
AND o.ORDER_TOTAL > 100000  
AND o.STATE = 'SHIPPED'  
AND c. CREATION_DATE > ?
```

## Used strategies to implement the query

### 1 + N strategy:

- Fetch recent customers
- Iterate through customers fetching their shipped orders
- Lots of round trips => **high-latency**

### Alternative strategy:

- Fetch recent customers
- Fetch recent orders
- Join
- 2 roundtrips but potentially large datasets => **inefficient**

**Not efficiently implemented using API Composition !**

# Solution (b) of challenge 2 : Implementing queries with CQRS

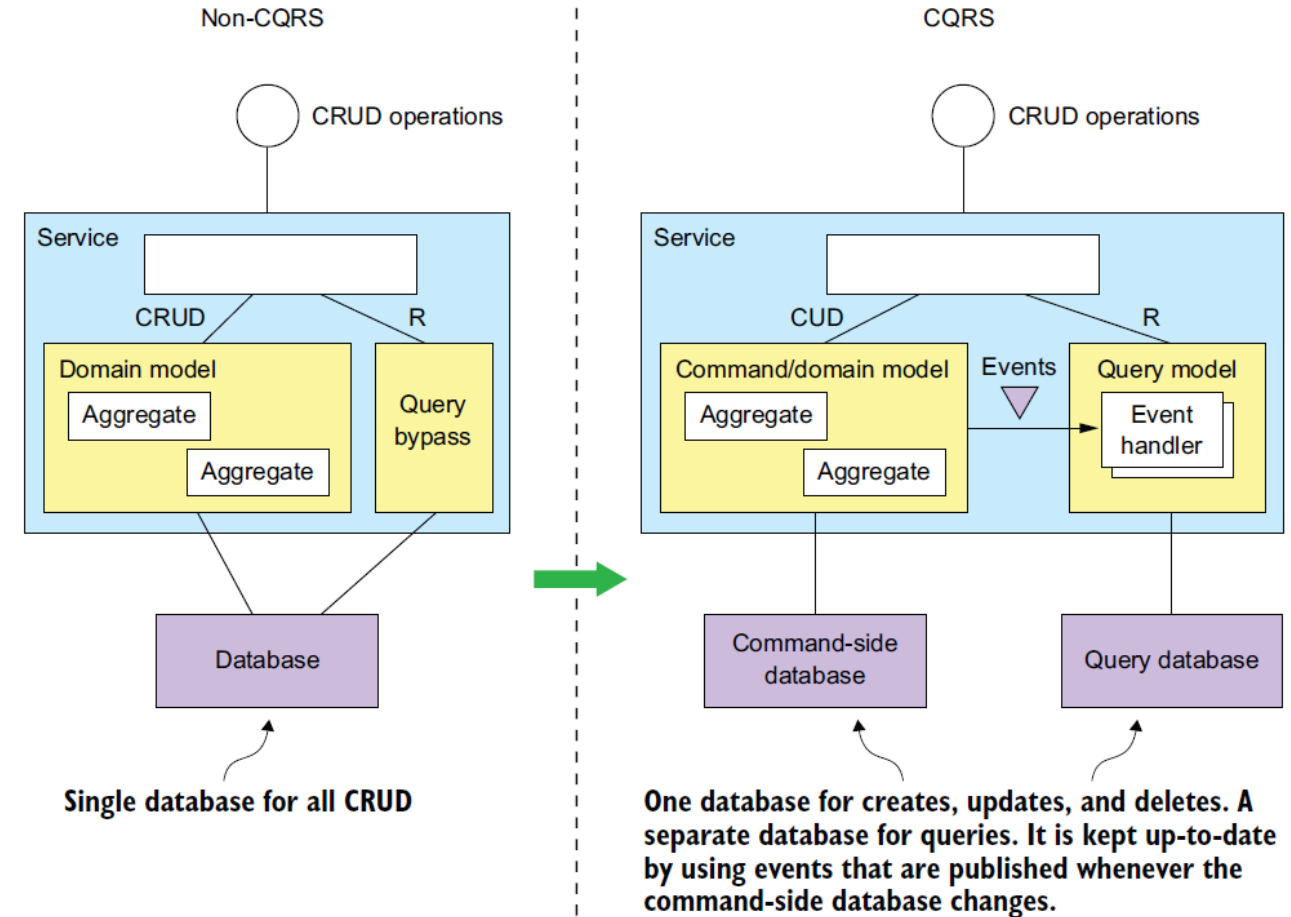
## CQRS (Command Query Responsibility Segregation)

**CQRS** = Using events to update a queryable replica

### The concept :

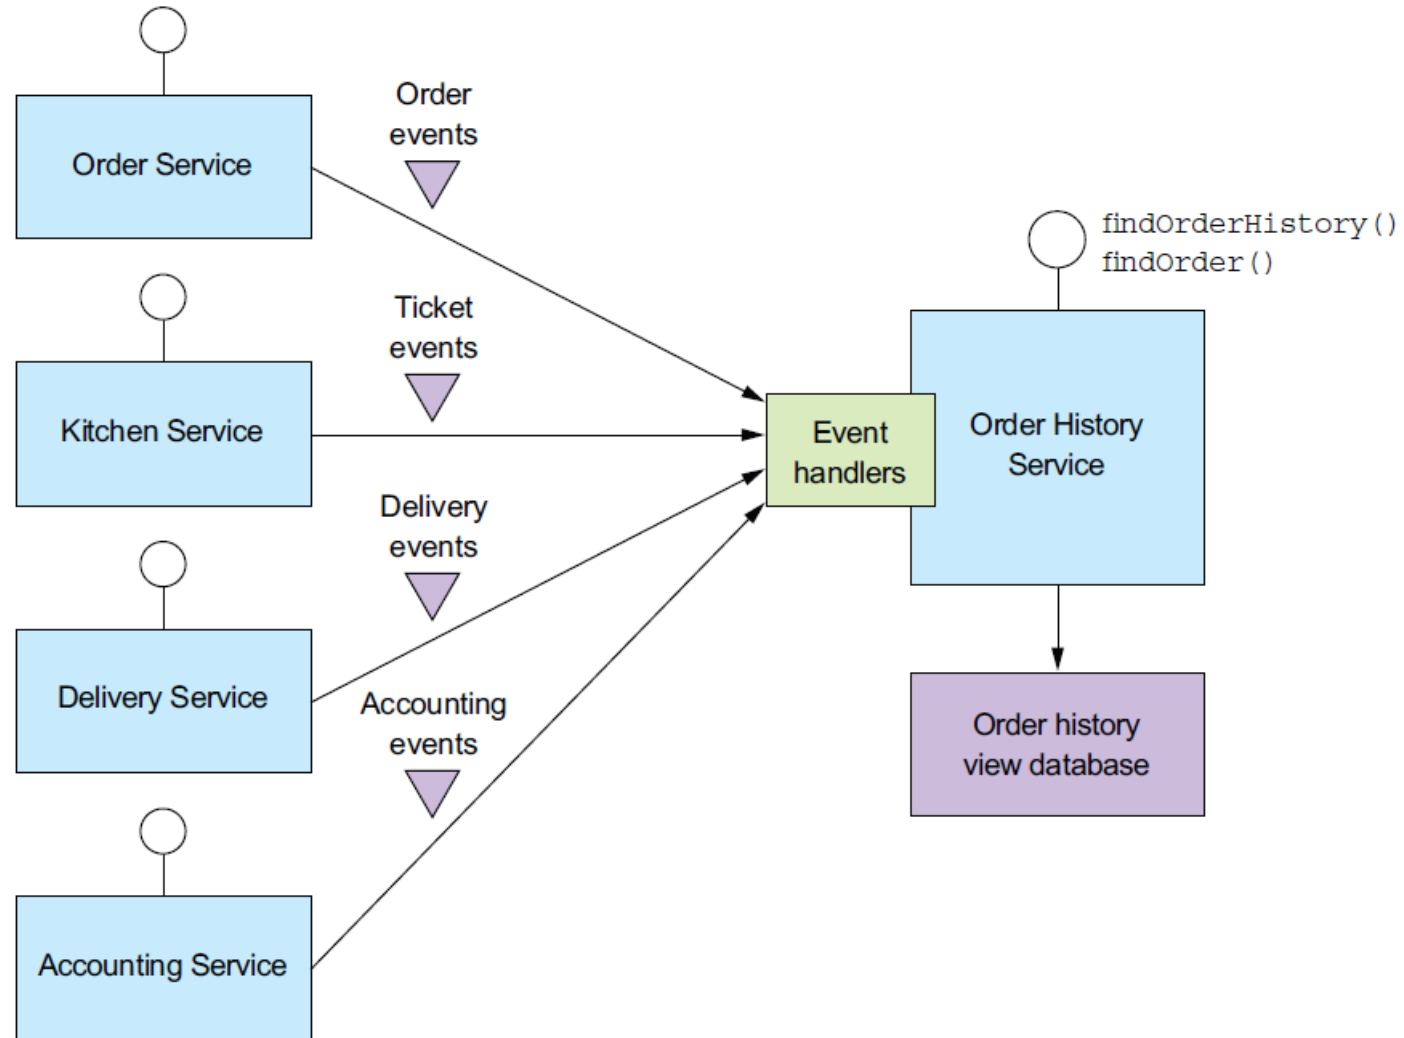
We define a **view database**, which is a **read-only replica** that is designed to support complex queries.

The application keeps the replica up to date by subscribing to events published by the services that own the data.



# Solution (b) of challenge 2 : Implementing queries with CQRS

## CQRS and Query-only services



# Solution (b) of challenge 2 : Implementing queries with CQRS

**View database :** Persisting a customer and order history in MongoDB

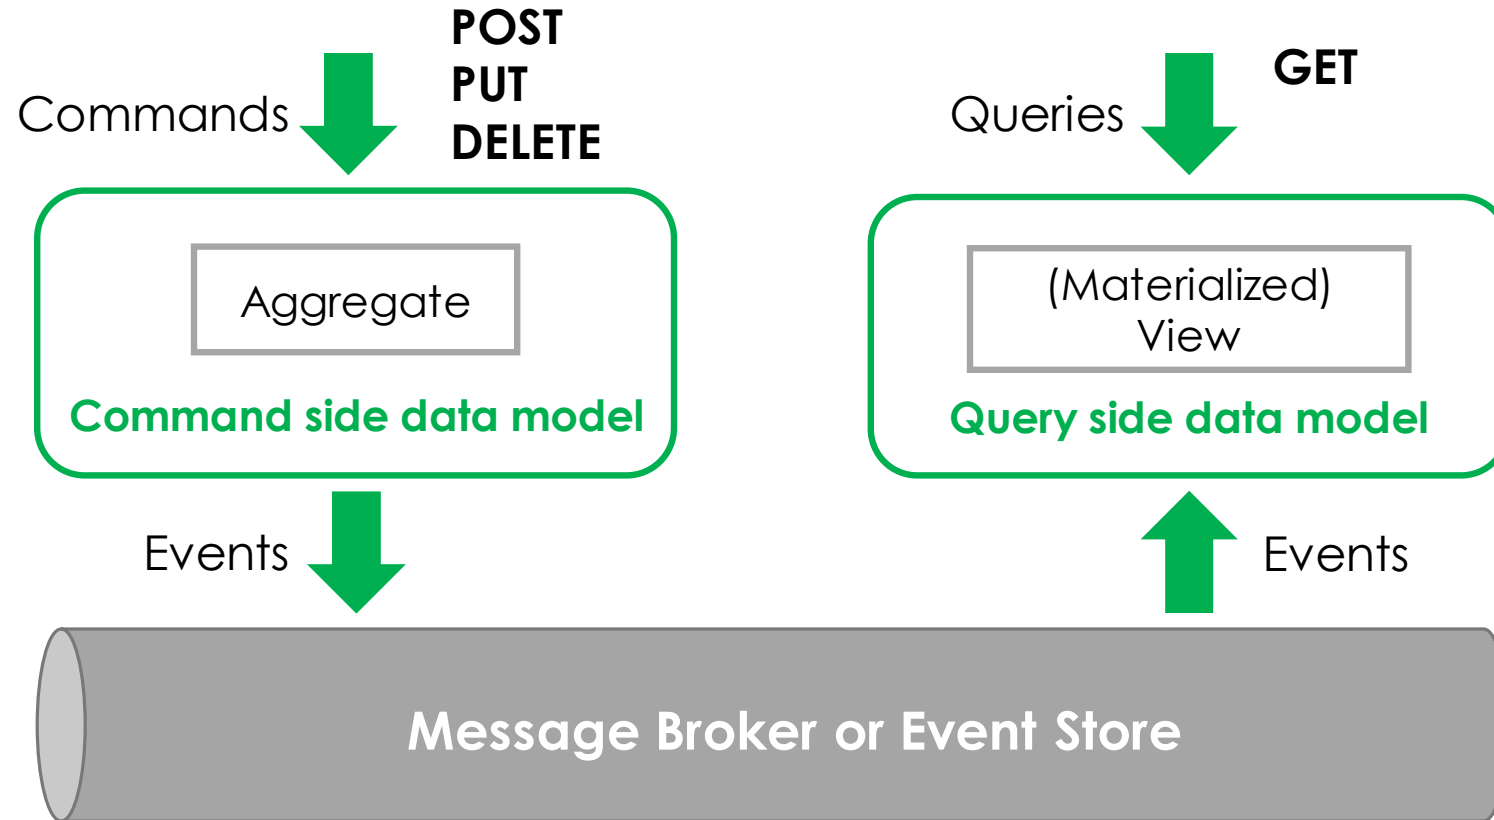
```
{
  "_id" : "0000014f9a45004b 0a0027000000000000",
  "name" : "Fred",
  "creditLimit" : {
    "amount" : "2000"
  },
  "orders" : {
    "0000014f9a450063 0a0027000000000000" : {
      "state" : "APPROVED",
      "orderId" : "0000014f9a450063 0a0027000000000000",
      "orderTotal" : {
        "amount" : "1234"
      }
    },
    "0000014f9a450063 0a0027000000000001" : {
      "state" : "REJECTED",
      "orderId" : "0000014f9a450063 0a0027000000000001",
      "orderTotal" : {
        "amount" : "3000"
      }
    }
  }
}
```

**Customer information**

**Order information**

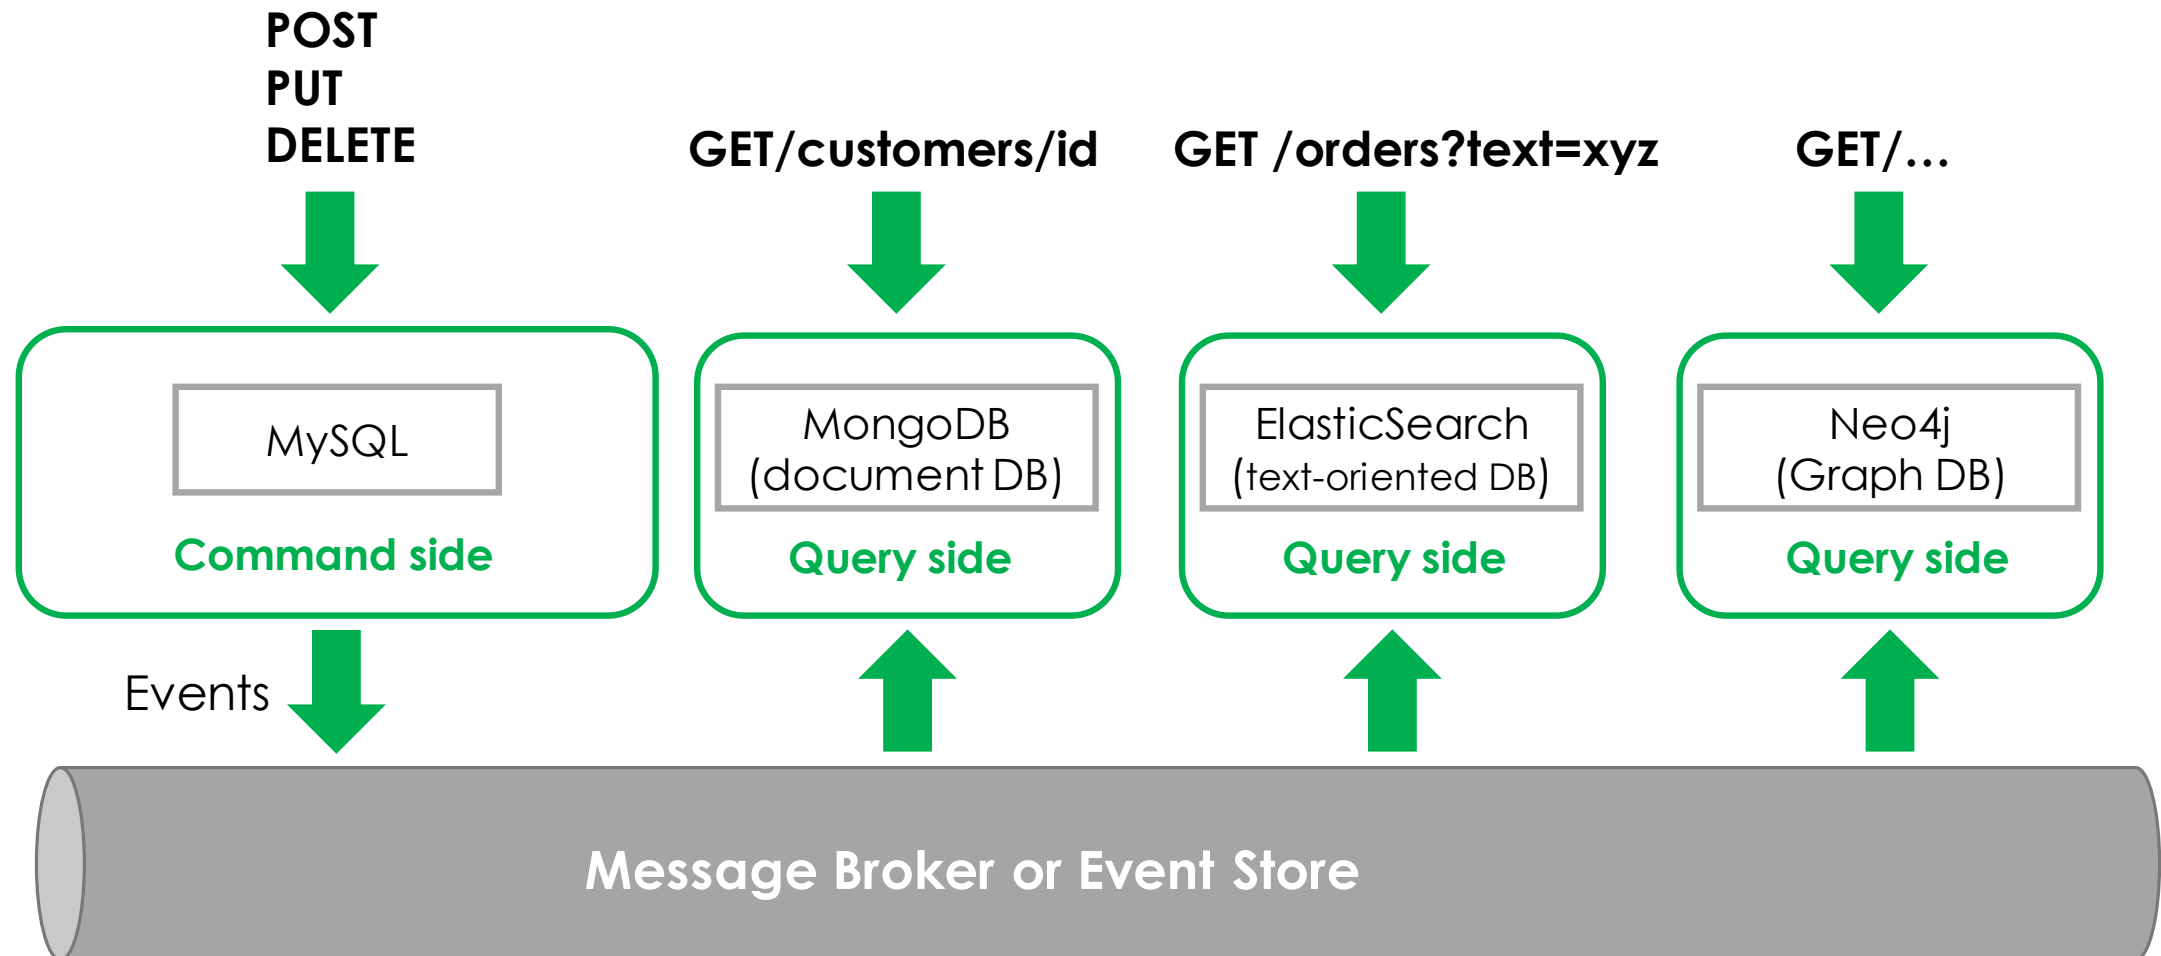
# Solution (b) of challenge 2 : Implementing queries with CQRS

Command Query Responsibility Segregation (CQRS)




# Solution (b) of challenge 2 : Implementing queries with CQRS

Queries with different database's types





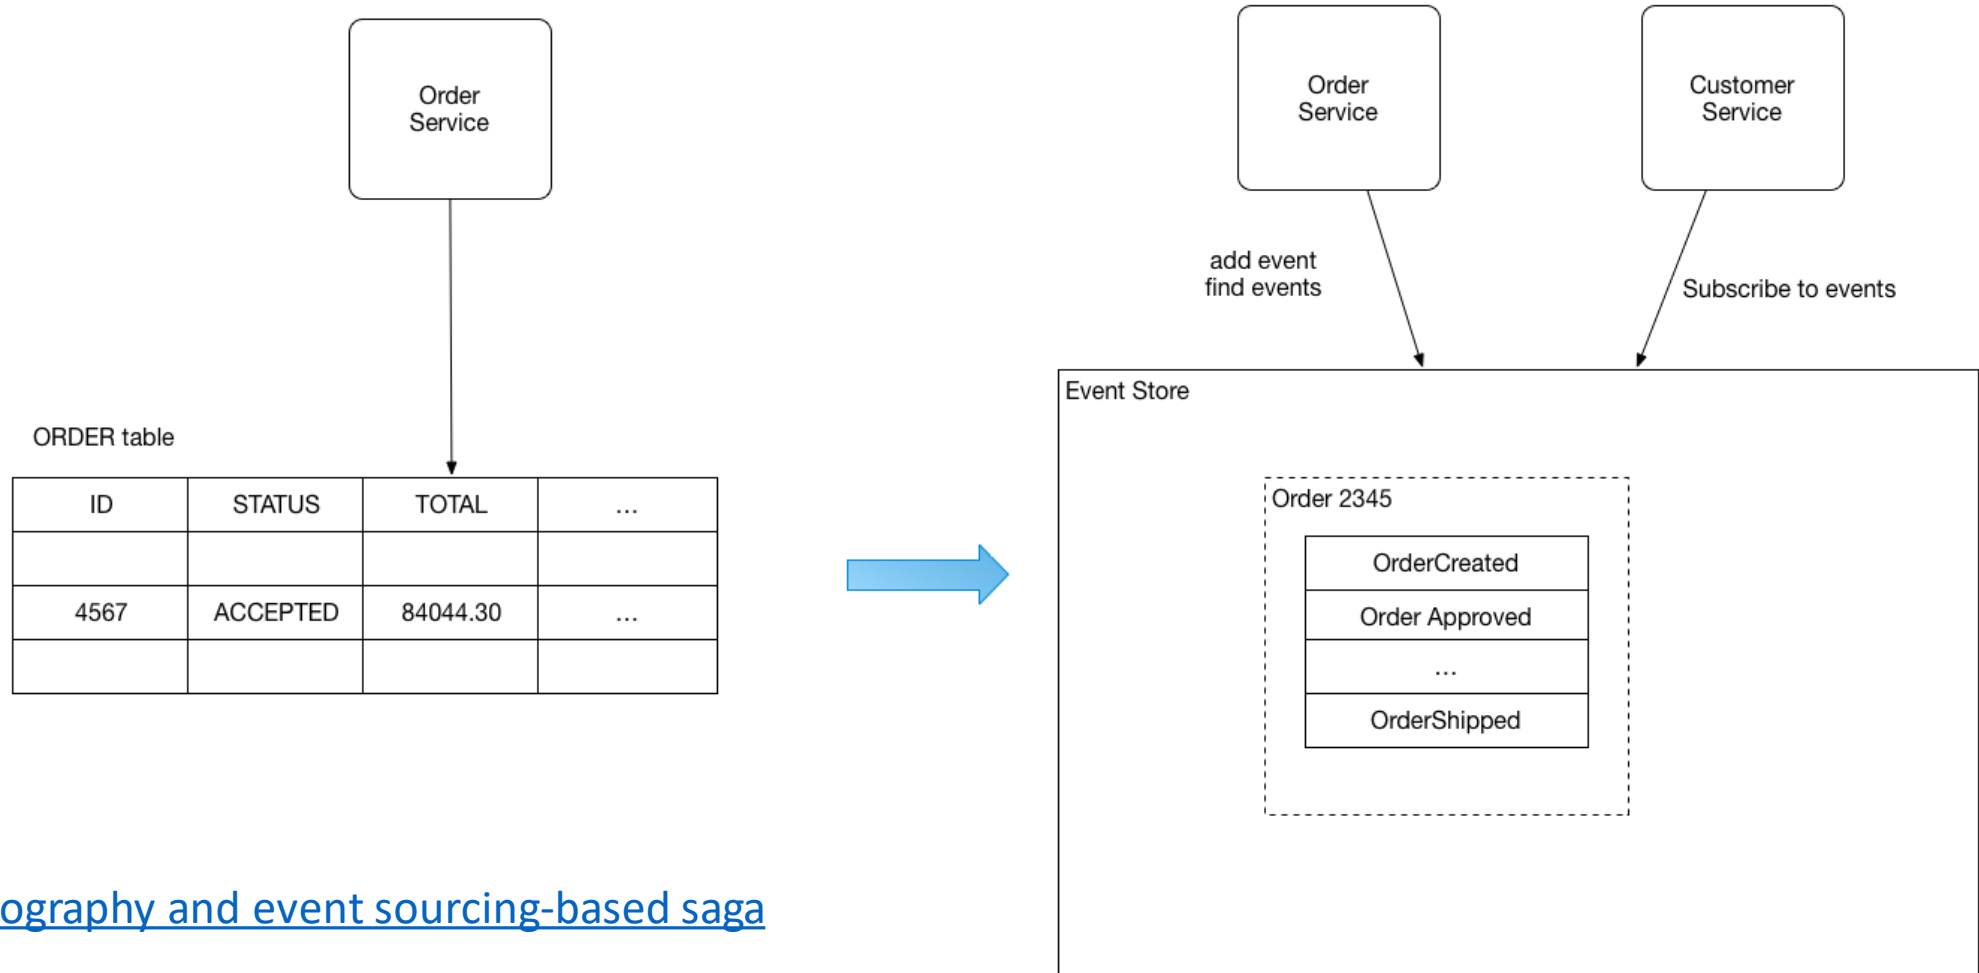


# Event-Driven Data Management

- Problem Statement : **Distributed Data Management** (*Transactions and Queries*)
- One solution of the **Transactions** challenge:
  - The **Saga** pattern
  - Coordination sagas
    - Choreography-based saga
    - Orchestration-based saga
- Two Solutions of the **Queries** challenge :
  - Solution 1 : The **API Composition** pattern
  - Solution 2 : The **CQRS** pattern
- BONUS: The **Event Sourcing** pattern

# Event sourcing

a.k.a. Persisting an object as a sequence of events




[Choreography and event sourcing-based saga](#)

# Event Sourcing

## (things to know)

- Event sourcing **persists the state of a business entity** such an *Order* or a *Customer* **as a sequence of state-changing events**.
- Whenever the state of a business entity changes, a new event is appended to the list of events.
- Since **saving an event** is a **single operation**, it is inherently **atomic**.
- The application reconstructs an entity's current state by replaying the events.
- Applications persist events in an **event store**, which is a **database of events**.
- The store has an API for adding and retrieving an entity's events.
- The event store also behaves like a message broker.
- It provides an API that enables services to subscribe to events.
- When a service saves an event in the event store, it is delivered to all interested subscribers.
- Some entities, such as a Customer, can have a large number of events. In order to optimize loading, an application can periodically save a snapshot of an entity's current state.
- To reconstruct the current state, the application finds the most recent snapshot and the events that have occurred since that snapshot.
- As a result, there are fewer events to replay.



# Event-Driven Data Management

</END>